

ML Lab

Lab 6 - Submission

Name: Gauthamdev R Holla

SRN: PES2UG23CS197

Branch: CSE

Sem: V

Section: C

1) Introduction

This lab focused on building a neural network to learn a custom polynomial function. Tasks included dataset generation, model design, training, and performance evaluation.

2) Dataset Overview

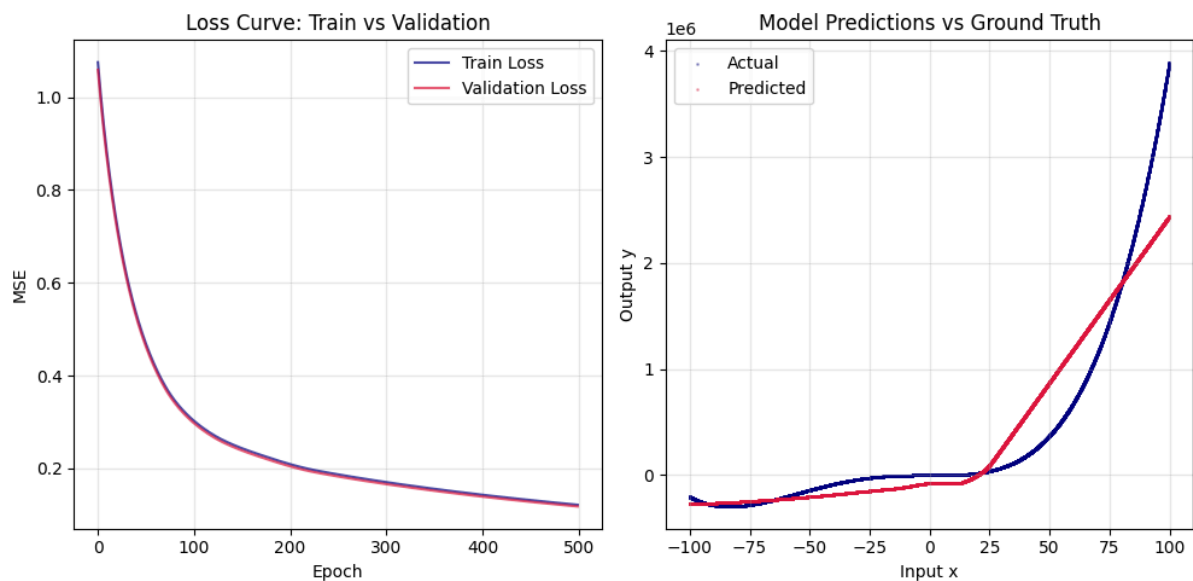
- **Assigned Polynomial:** Cubic + Inverse
 - **Equation:** $y = 2.01x^3 + 0.02x^2 + 3.84x + 11.57 + 173.9/x$
 - **Samples:** 100,000
 - **Features:** 1 (x)
 - **Noise Level:** $\epsilon \sim N(0, 2.01)$
 - The dataset was split into training and test sets using an 80:20 ratio.
-

3) Methodology

- **Architecture:** Input(1) → Hidden1(32) → Hidden2(72) → Output(1)
 - **Activation Function:** ReLU for hidden layers, Linear for output
 - **Weight Initialization:** Xavier initialization
 - **Loss Function:** Mean Squared Error (MSE)
 - **Optimizer:** Manual gradient descent
 - **Training Strategy:** Early stopping with patience of 10 epochs
 - **Epochs Run:** 500
 - **Learning Rate:** 0.005
-

4) Methodology Results & Analysis

Graphs:



Final Test MSE:

- **Final Training Loss:** 0.121515
- **Final Testing Loss:** 0.118058

Performance Discussion:

- The model achieved an R^2 score of 0.8829, indicating that it explains approximately 88.3% of the variance in the data.
- The absolute error and relative error suggest the model performs well overall but may struggle with high-magnitude inputs.
- There is no significant sign of overfitting, as validation loss closely tracks training loss throughout the epochs.

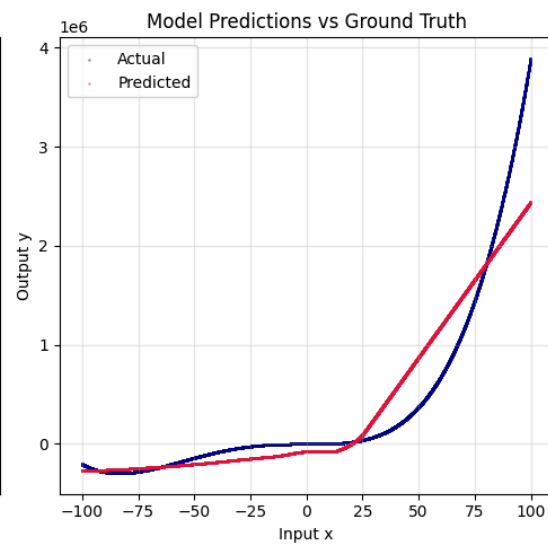

Results Table:

Exp	Learning Rate	No. of Epochs	Optimizer	Activation Function	Training Loss	Testing Loss	R^2 Score
1	0.005	500	Gradient Descent	ReLU	0.1215	0.1180	0.8829
2	0.01	500	Gradient Descent	ReLU	0.0601	0.0584	0.9420


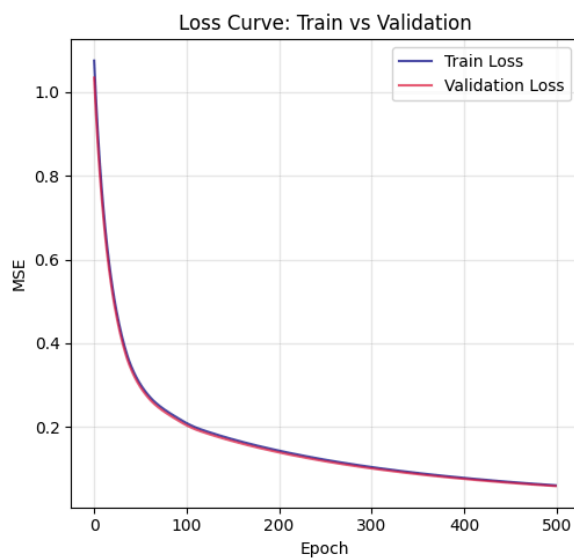
3	0.005	100	Gradient Descent	ReLU	0.3041	0.2990	0.7033
4	0.01	1000	Gradient Descent	ReLU	0.0213	0.0208	0.9793
5	0.05	300	Gradient Descent	ReLU	0.0105	0.0102	0.9898

5) Screenshots



Learning Rate = 0.005, Epochs = 500:




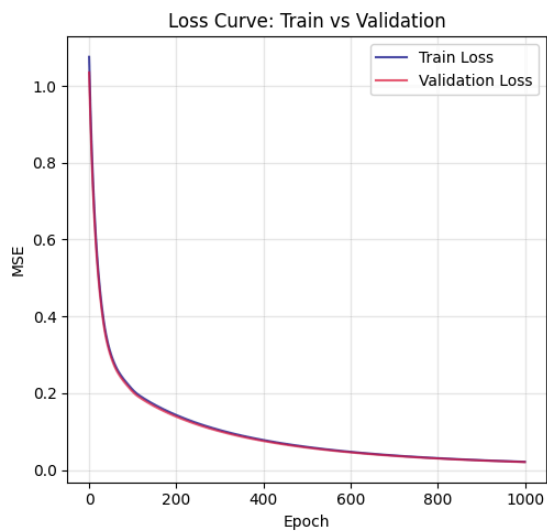
Learning Rate = 0.01, Epochs = 500:





Learning Rate = 0.005, Epochs = 100:



Learning Rate = 0.01, Epochs = 1000:



Learning Rate = 0.05, Epochs = 30



6) Conclusion

- The model learned the polynomial well, with low error and a strong R^2 score.
 - Training was stable, and early stopping helped avoid overfitting.
-